Extending SDN to the Data Plane

Anirudh Sivaraman, Keith Winstein, Suvinay Subramanian,
Hari Balakrishnan

M.I.T.
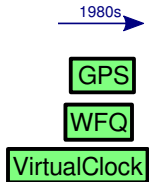
http://web.mit.edu/anirudh/www/sdn-data-plane.html
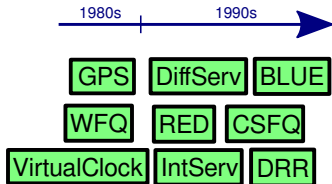
November 7, 2013

Switch Data Planes today

Two key decisions on a per-packet basis:

- Scheduling: Which packet to transmit next?

- Queue Management: How long can queues grow? Which packet to drop?

The long lineage of in-network algorithms
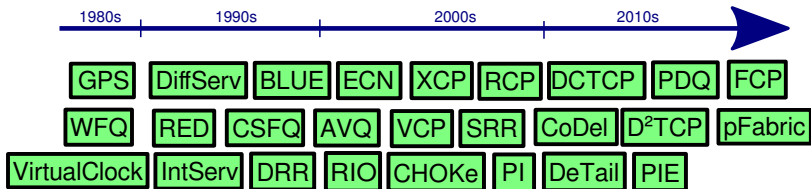
1980s

GPS

WFQ

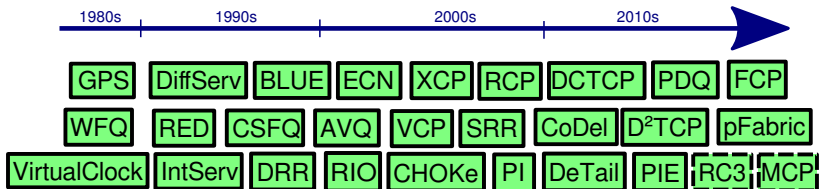VirtualClock

The long lineage of in-network algorithms

The long lineage of in-network algorithms

The long lineage of in-network algorithms

The long lineage of in-network algorithms

The Data Plane is continuously evolving

- Each scheme wins in its own evaluation.

- Some believe in a "silver bullet" knobless in-network method.

We disagree: There is no silver bullet!

- Different applications care about different objectives.

- Applications use different transport protocols.

- Networks are heterogeneous.

Our work:

- Quantify non-universality of in-network methods.

- Extend SDN to the Data Plane.

Quantifying "No Silver Bullet": Network Configurations

Configuration	Description
CoDel+FCFS	One shared FCFS queue with CoDel
CoDel+FQ	Per-flow fair queueing with CoDel on each queue
Bufferbloat+FQ	Per-flow fair queueing with deep buffers on each queue

Quantifying "No Silver Bullet": Workloads and Objectives

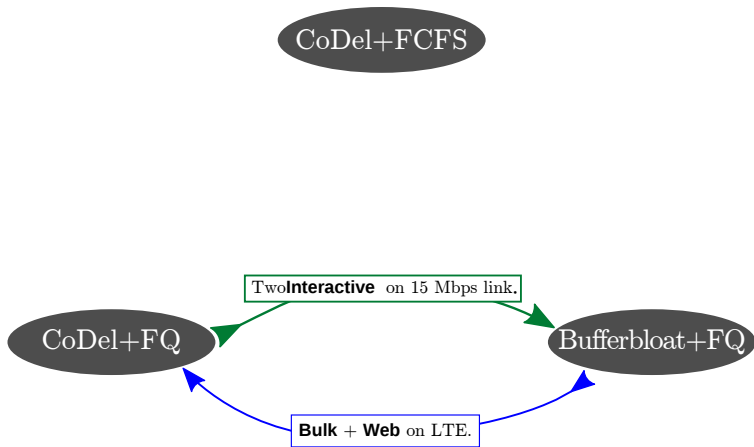| Workload	Description	Objective
Bulk	Long-running bulk transfer flow	Max. throughput
Web	Switched flow with ON and OFF periods	Min. 99.9 %ile flow completion time
Interactive	Long-running interactive flow	Max. $\dfrac{\text{throughput}}{\text{delay}}$ (power)

Quantifying "No Silver Bullet"
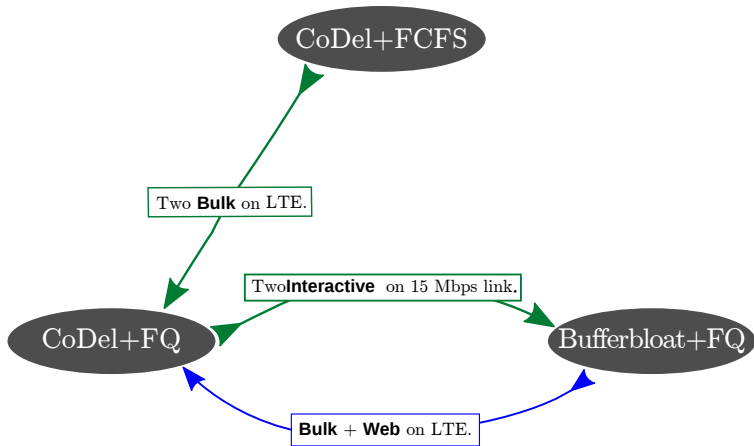
CoDel+FCFS

CoDel+FQ

Bufferbloat+FQ

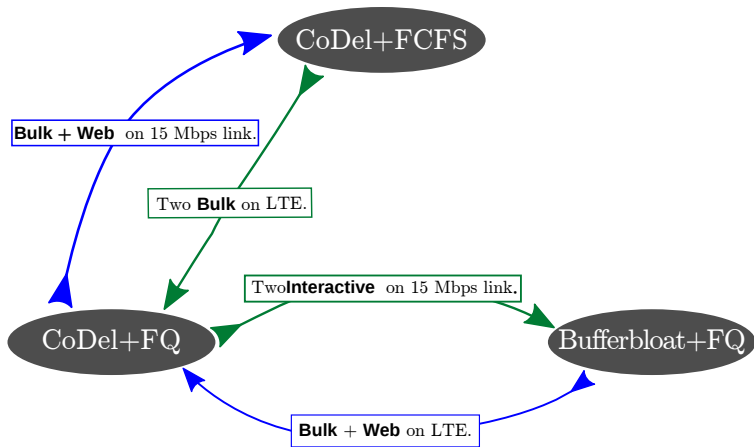# Quantifying "No Silver Bullet"

Quantifying "No Silver Bullet"

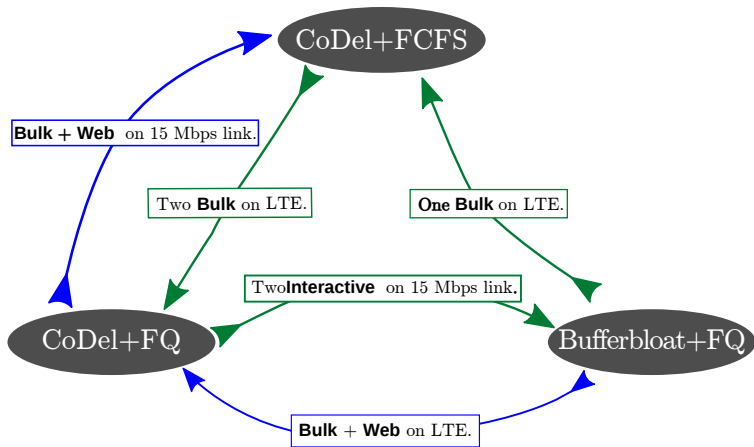# Quantifying "No Silver Bullet"

Quantifying "No Silver Bullet"

Quantifying "No Silver Bullet"

Quantifying "No Silver Bullet"
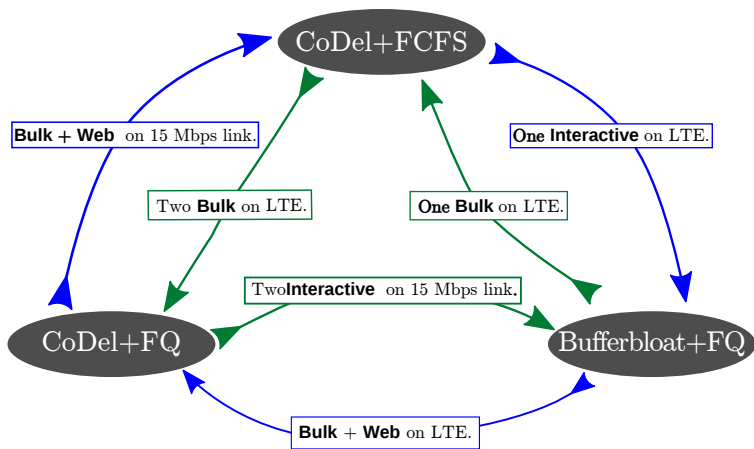
Why is no single data plane configuration the best?

- Bufferbloat on variable-rate links helps throughput!
  - Variable-rate links have an inherent delay-throughput tradeoff

- FCFS is preferable to Fair Queuing in some cases
  - When equally aggressive flows compete, they don't need protection from each other
  - Helps reduce tail packet delay

- Fair Queuing is required in some cases
  - When competing flows aren't equally aggressive, isolation helps

So what should the network designer do?

Architect a flexible data plane

- ▶ Programmable queue management and scheduling
- ▶ Not just for selecting among pre-built choices, but to change behavior in the field
- ▶ Because there is no silver bullet and innovation will continue!

Controlled flexibility: Want performance, security

- ► Provide interfaces only to the head and tail of queues
- ► Operators specify only queue-management/scheduling logic
- ► No access to packet payloads.

Building such a data plane in four parts

- Hardware gadgets
  - Random number generators (RED, BLUE)
  - Binary tree of comparators (pFabric, SRPT)
- I/O interfaces
  - Drop/mark head/tail of queue
  - Interrupts for enqueue/dequeue
- State maintenance
  - Per-flow (WFQ, DRR)
  - Per-dst address (PF)
- A domain-specific instruction set
  - Expresses control flow
  - Implements new functions unavailable in hardware

Feasibility study: CoDel

Synthesis numbers on Xilinx Kintex-7:

Resource	Usage	Fraction
Slice logic	1,256	1%
Slice logic dist.	1,975	2%
IO/GTX ports	27	2%
DSP slices	0	0%
Maximum speed	12.9×10^6 pkts/s ~10gbps	

- Small fraction of the FPGA's resources.
- Can be improved by pipelining or parallelizing.

Limitations and Practical Considerations:

- Cannot express several network functions that need payloads.
- How do applications signal objectives to the network?
- Feasibility at 10G on high port-density switches.
- Energy and Area overheads.

Related Work:

- Active Networking, e.g., ANTS

- Software Routers, e.g., Click

- Software-Defined Networking, e.g., OpenFlow

Conclusion

- There is no silver bullet to in-network resource control because of application and network diversity

- Algorithms will continue to evolve: the data plane should help

- Directions to reproduce results: http://web.mit.edu/anirudh/www/sdn-data-plane.html