Extending SDN to the Data Plane

Anirudh Sivaraman, Keith Winstein, Suvinay Subramanian,
Hari Balakrishnan

M.I.T.

http://web.mit.edu/anirudh/www/sdn-data-plane.html

Switch Data Planes today

Two key decisions on a per-packet basis:

- **Scheduling**: Which packet to transmit next?
- **Queue Management**: How long can queues grow? Which packet to drop?

The long lineage of in-network algorithms

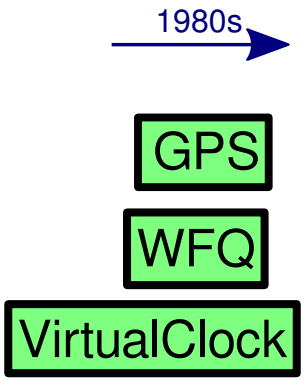

The long lineage of in-network algorithms

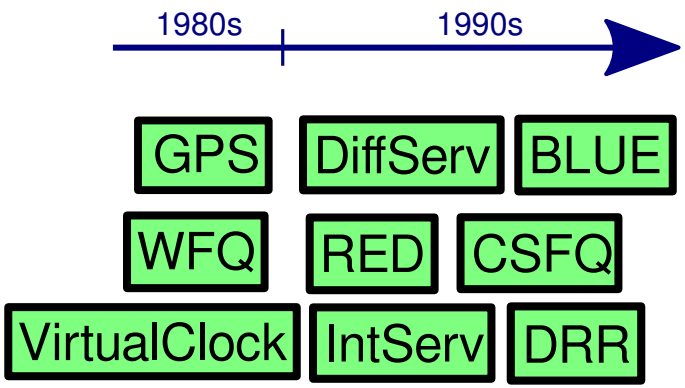

The long lineage of in-network algorithms

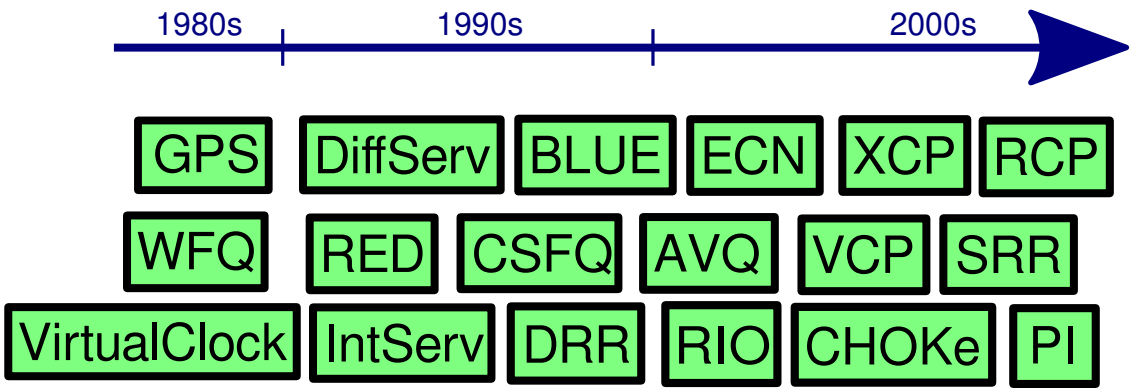

The long lineage of in-network algorithms

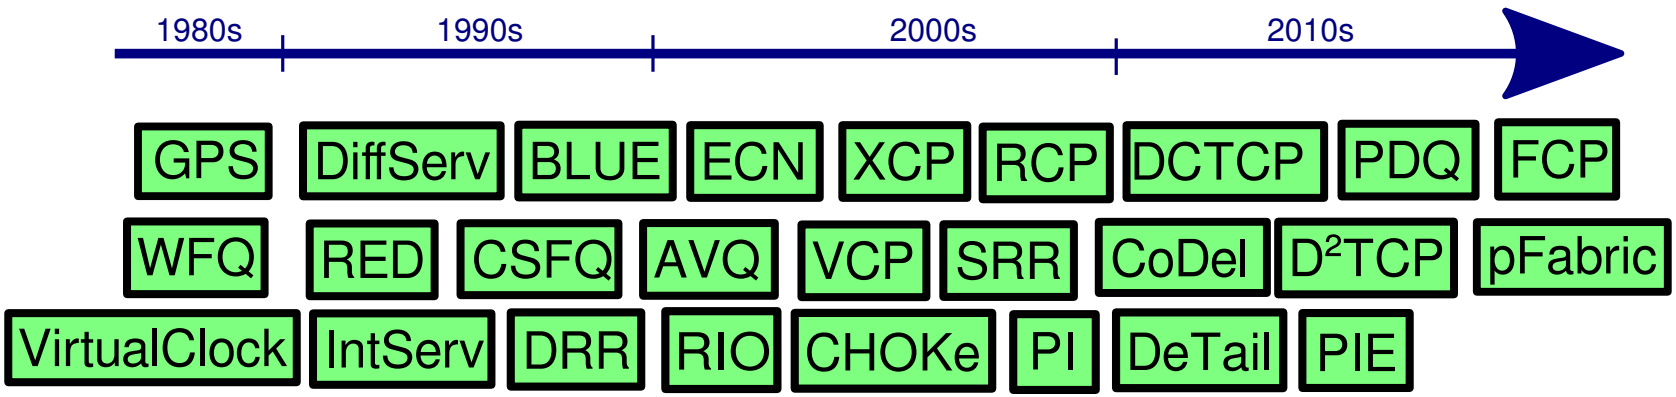

The long lineage of in-network algorithms

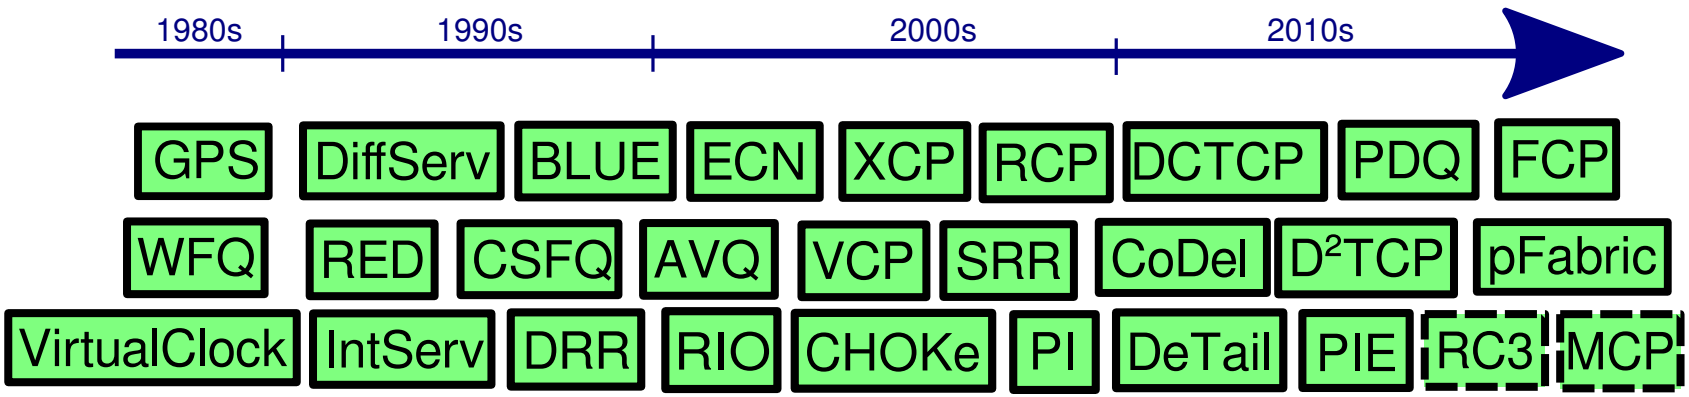

The Data Plane is continuously evolving

- Each scheme wins in its own evaluation.
- Quest for a “silver bullet” in-network method.

We disagree: There is no silver bullet!

- Different applications care about different objectives.
- Applications use different transport protocols.
- Networks are heterogeneous.

Our work:

- Quantify non-universality of in-network methods.
- Extend SDN to the Data Plane to handle in-network diversity.

Quantifying "No Silver Bullet": Network Configurations

Configuration	**Description**
CoDel+FCFS	One shared FCFS queue with CoDel
CoDel+FQ	Per-flow fair queueing with CoDel on each queue (Nichols 2013)
Bufferbloat+FQ	Per-flow fair queueing with deep buffers on each queue

Quantifying "No Silver Bullet": Workloads and Objectives

Workload	**Description**	**Objective**
Bulk	Long-running bulk transfer flow	Max. throughput
Web	Switched flow with ON/OFF periods	Min. 99.9 %ile flow completion time
Interactive	Long-running interactive flow	Max. $\frac{\text{throughput}}{\text{delay}}$

Quantifying "No Silver Bullet"

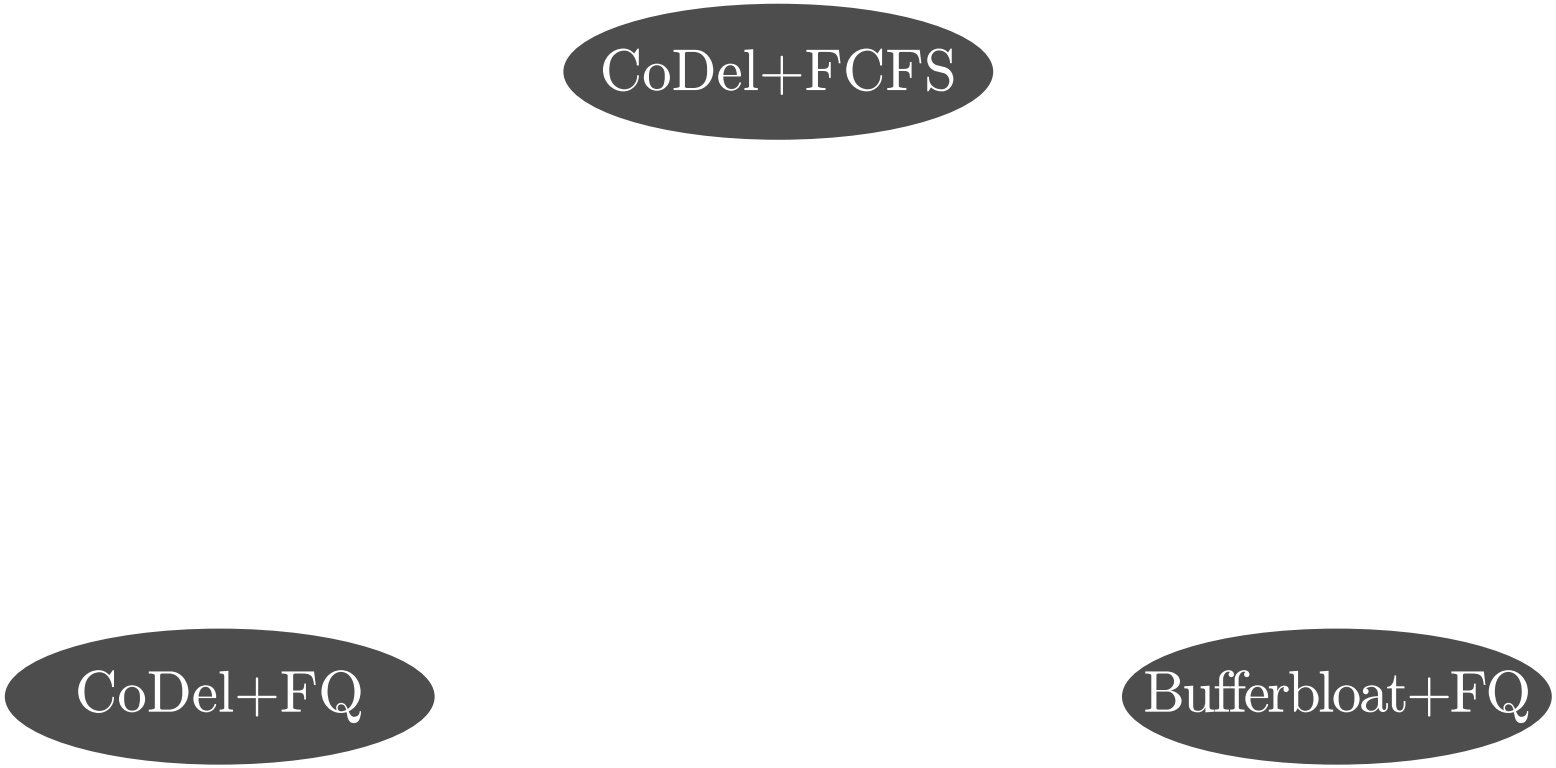

Quantifying "No Silver Bullet"

Quantifying "No Silver Bullet"

CoDel+FCFS

Experiment configuration:
Workload: 1 Bulk flow + 1 Web Flow
Network: LTE link with 150 ms min. RTT

CoDel+FQ

Bufferbloat+FQ

Quantifying "No Silver Bullet"

CoDel+FCFS

Experiment configuration:
Workload: 1 Bulk flow + 1 Web Flow
Network: LTE link with 150 ms min. RTT

Bulk Tpt: 3.9 Mbps

CoDel+FQ

Bufferbloat+FQ

Web Tail FCT: 43 s

Quantifying "No Silver Bullet"

CoDel+FCFS

Experiment configuration:
Workload: 1 Bulk flow + 1 Web Flow
Network: LTE link with 150 ms min. RTT

Bulk Tpt: 3.9 Mbps

CoDel+FQ

Web Tail FCT: 43 s

Bulk Tpt: 11.2 Mbps

Bufferbloat+FQ

Web Tail FCT: 21 s

Quantifying "No Silver Bullet"

CoDel+FCFS

Experiment configuration:
Workload: 1 Bulk flow + 1 Web Flow
Network: LTE link with 150 ms min. RTT

Bulk Tpt: 3.9 Mbps

CoDel+FQ

Web Tail FCT: 43 s

Bulk Tpt: 11.2 Mbps

Bufferbloat+FQ

Web Tail FCT: 21 s

Quantifying "No Silver Bullet"

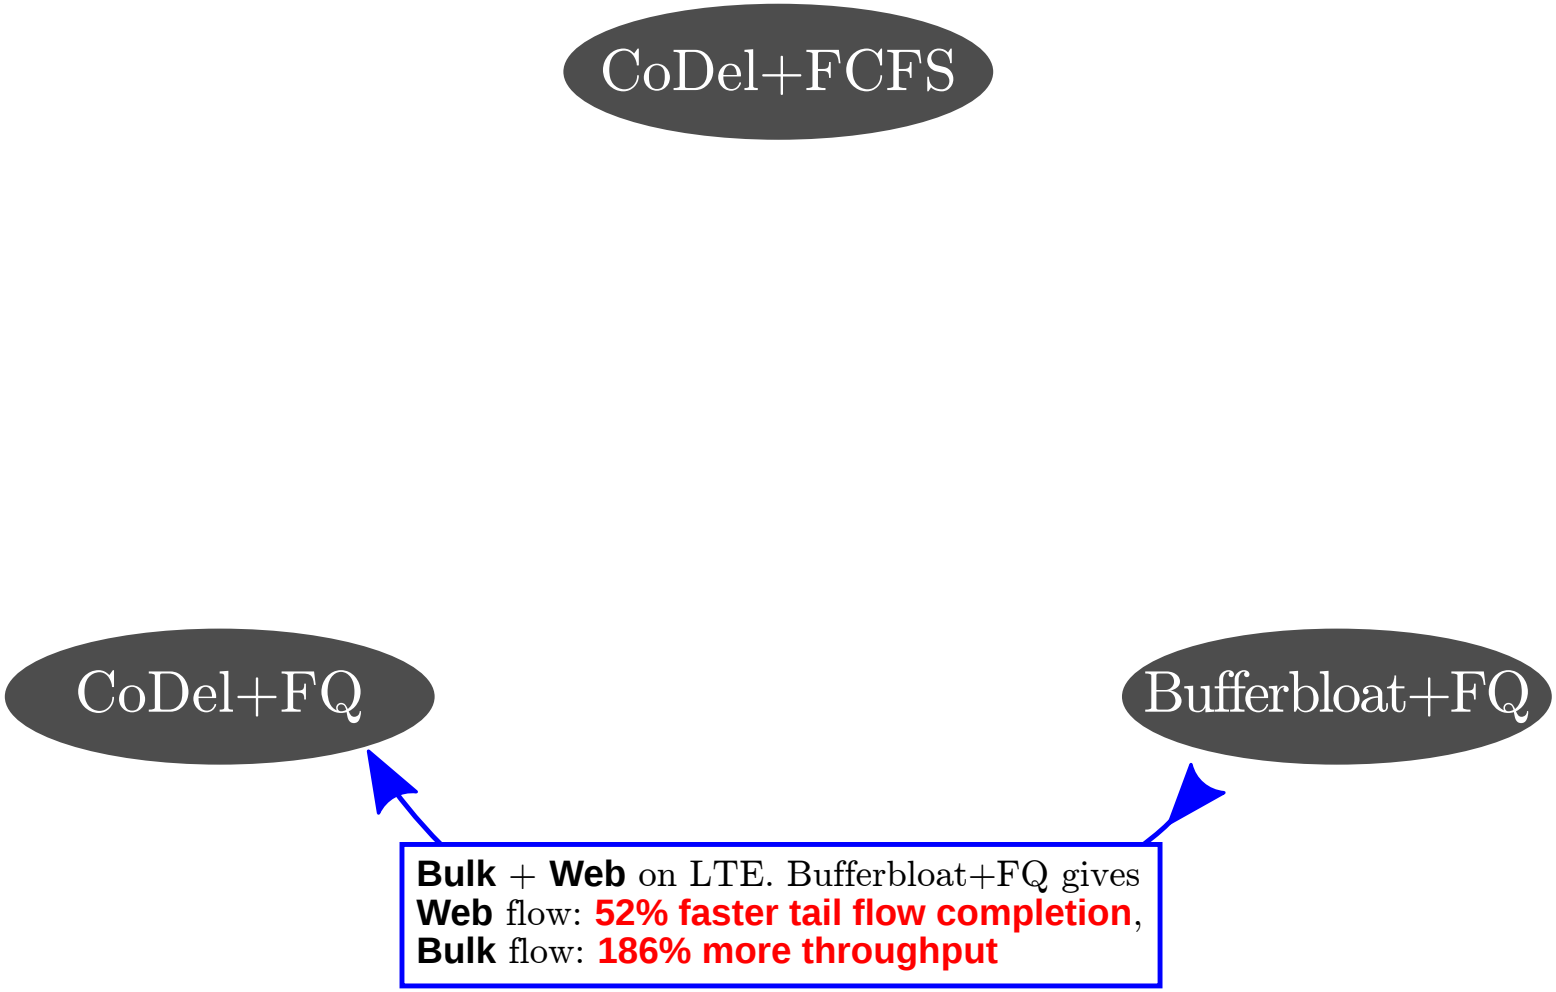

Quantifying "No Silver Bullet"

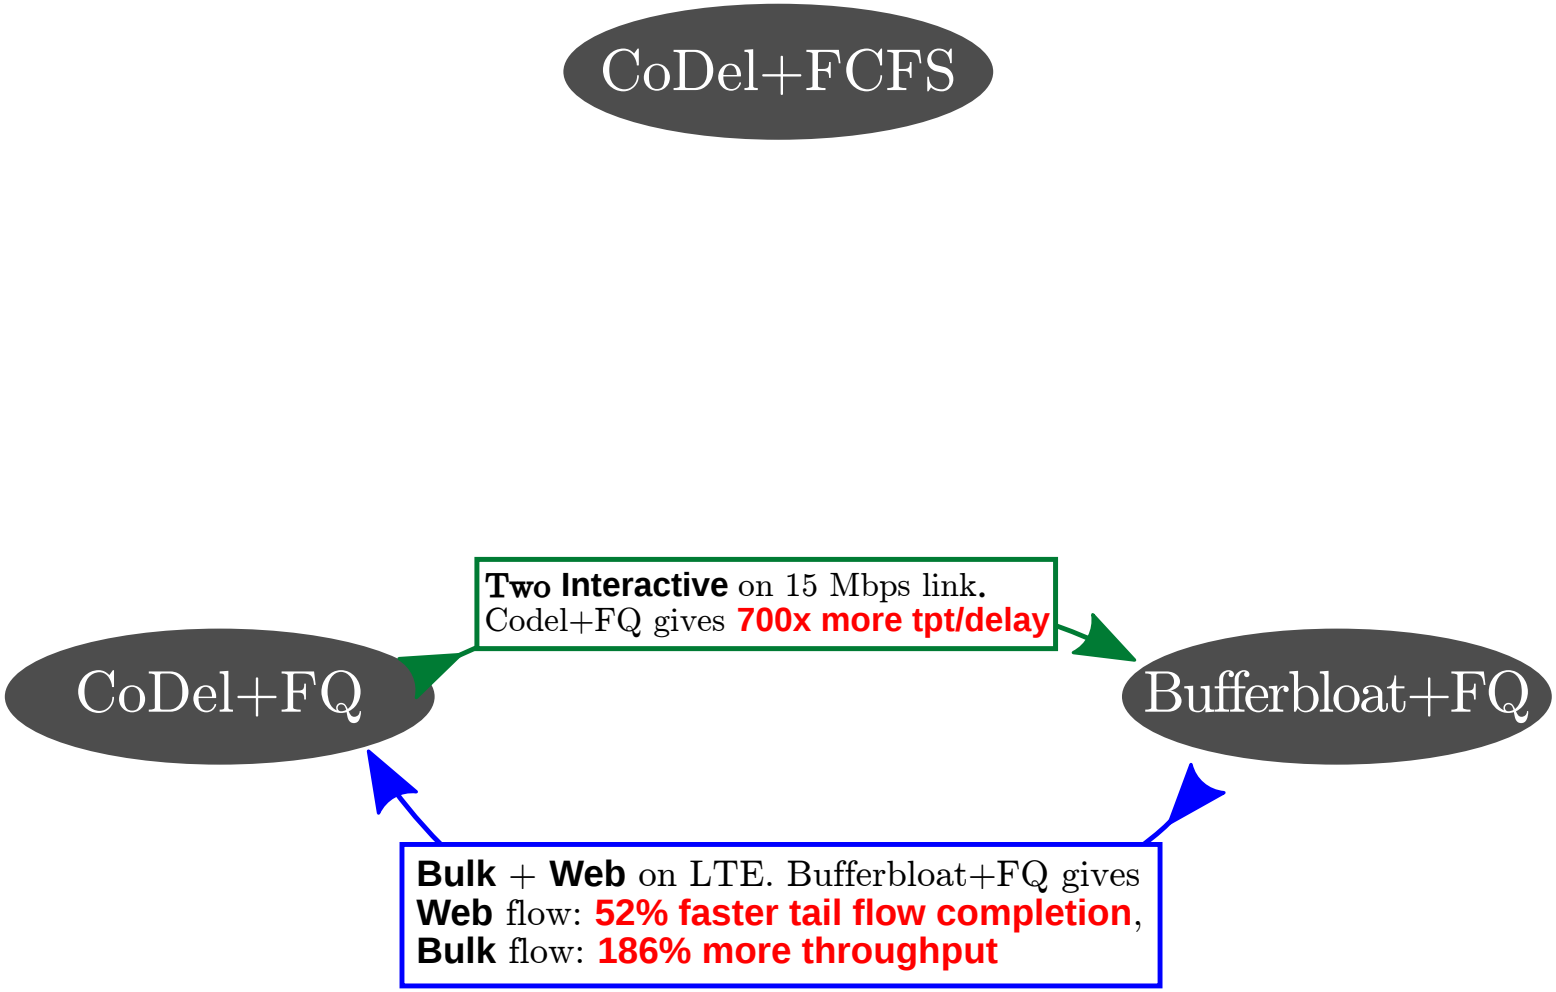

Quantifying "No Silver Bullet"

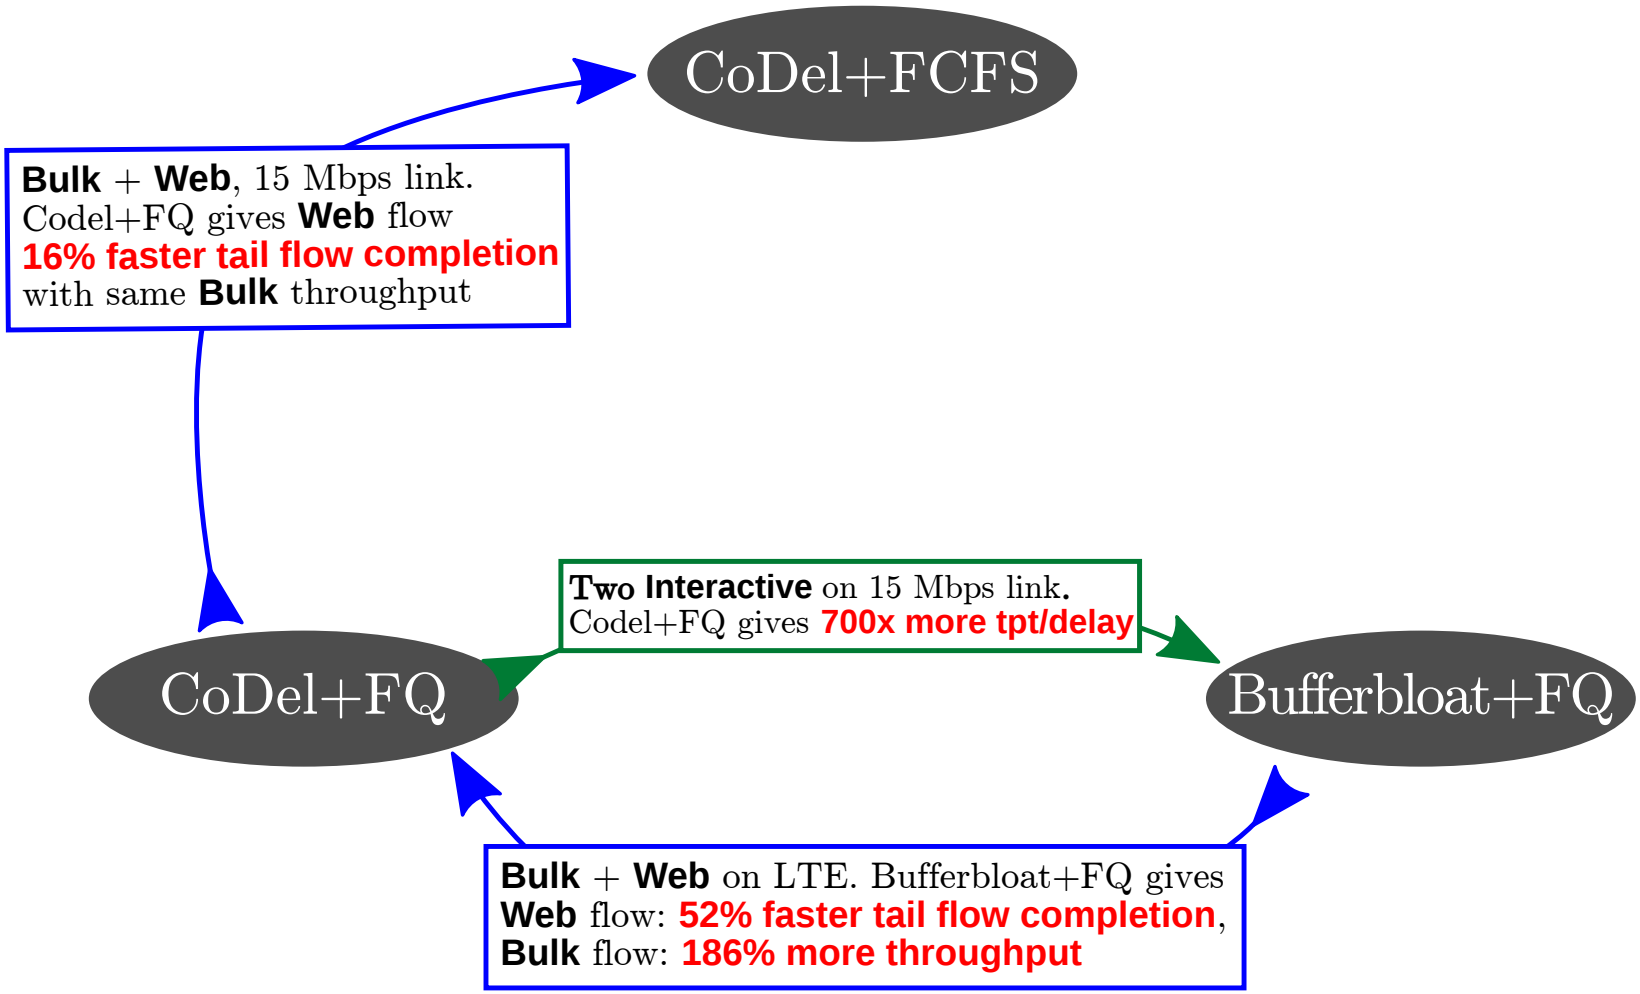

Quantifying "No Silver Bullet"

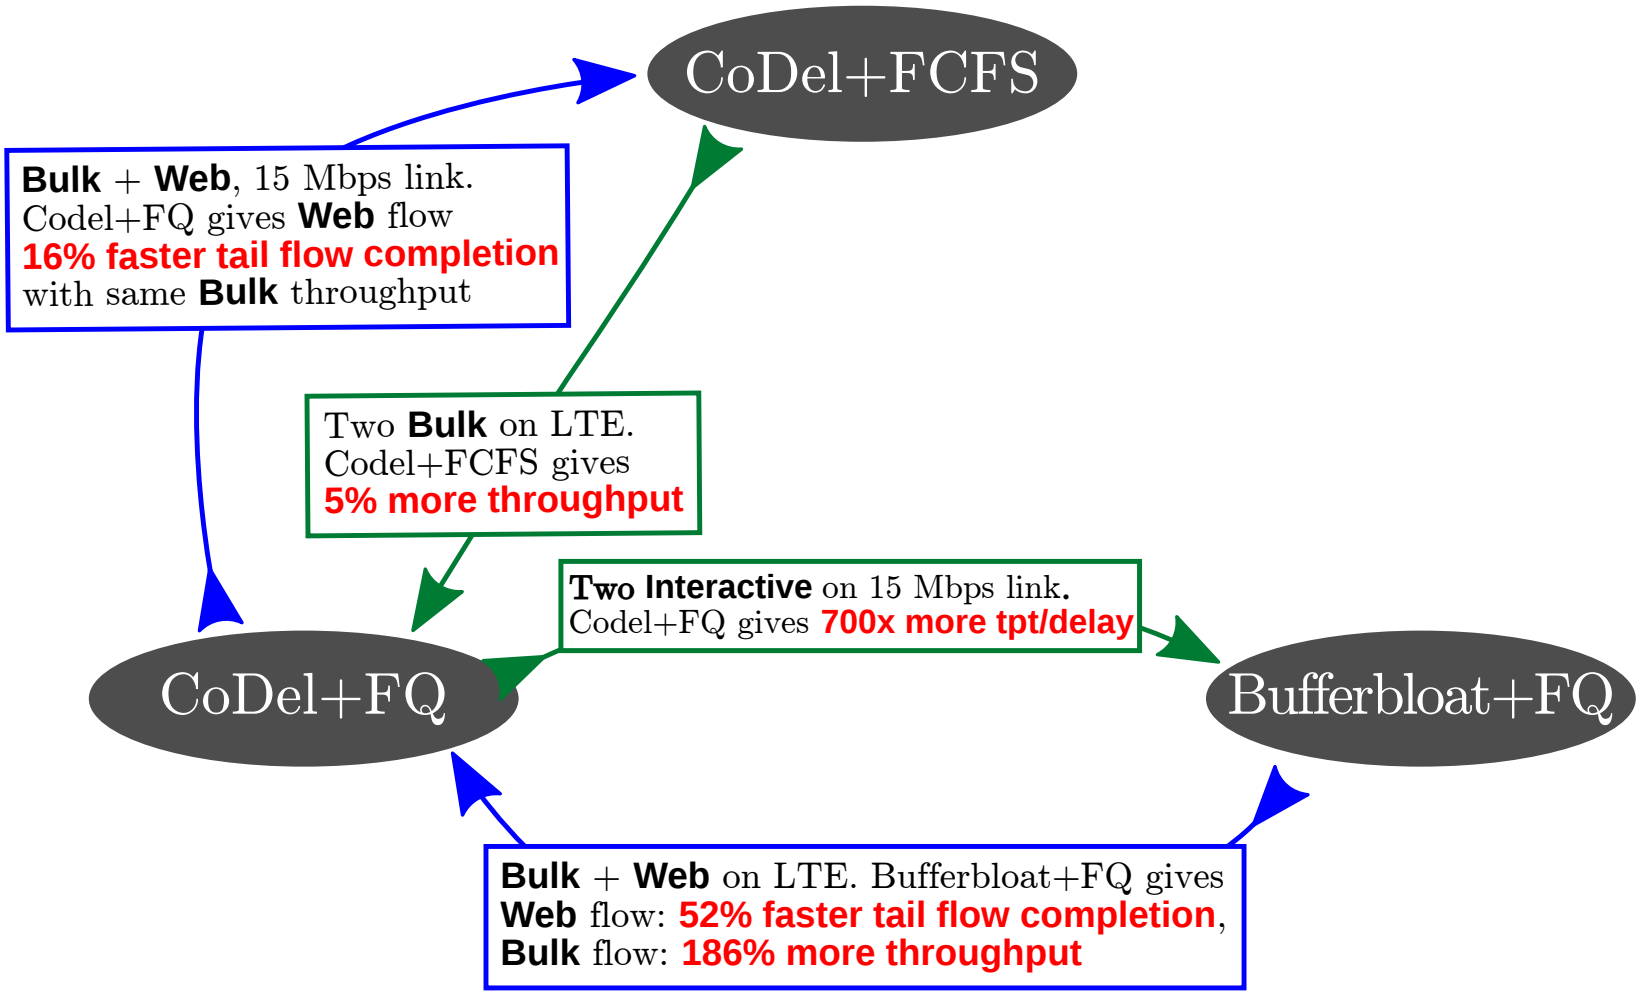

Quantifying "No Silver Bullet"

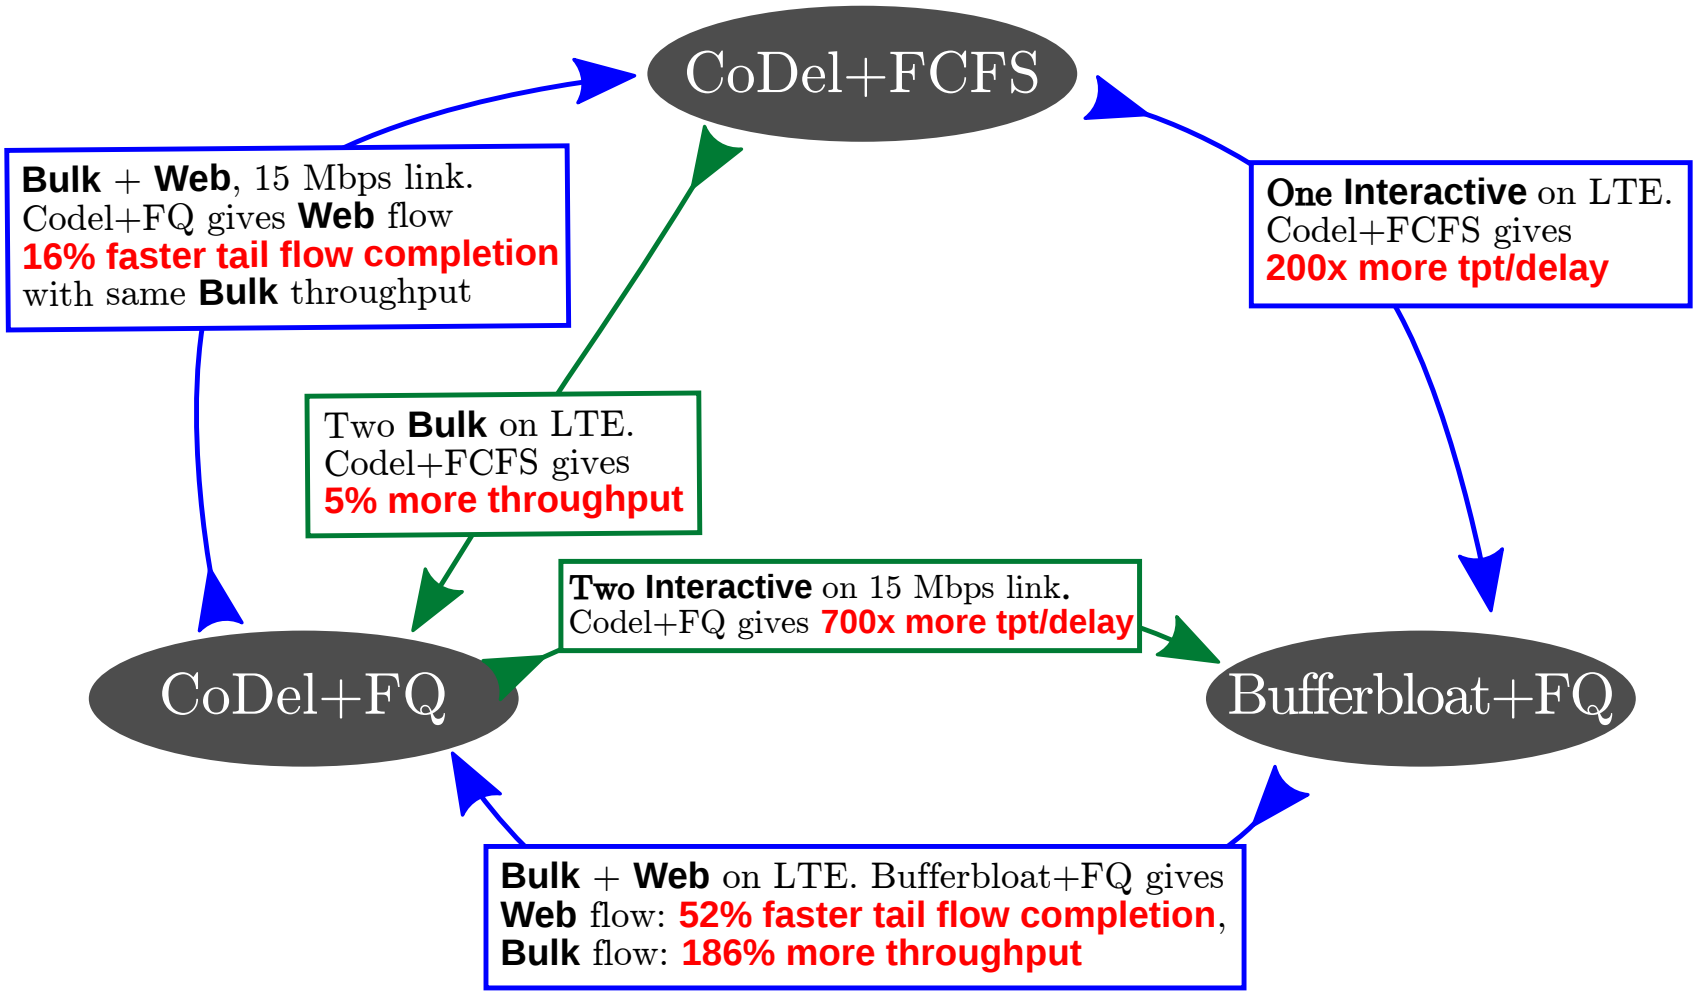

Quantifying "No Silver Bullet"

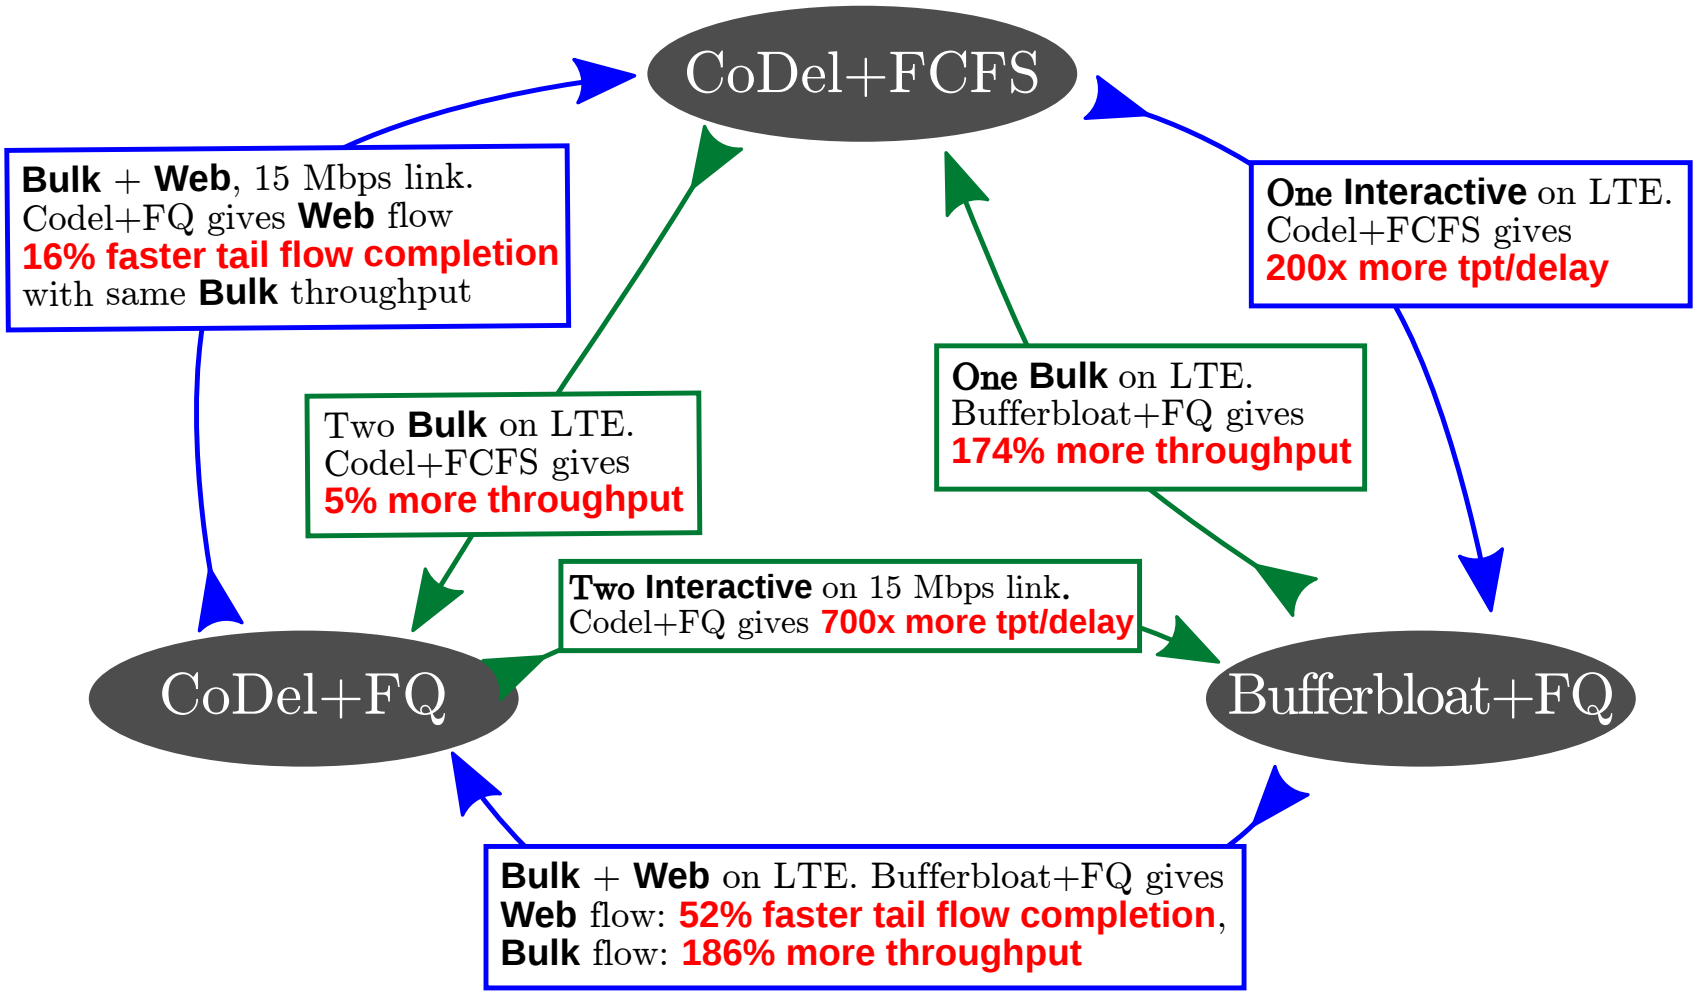

Why is no single data plane configuration the best?

- Bufferbloat gives the best throughput on variable-rate links.
- FCFS is preferable to Fair Queuing with homogenous objectives.
- Fair Queuing is preferable with heterogeneous objectives.

So what should the network designer do?

- Don't strive for the best in-network behaviour.
- Instead, architect for evolvability.
- Conceptually, extend SDN to include the data plane as well.

Flexibility without sacrificing performance

- Provide interfaces only to the head and tail of queues
- Operators specify only queue-management/scheduling logic
- No access to packet payloads.

Building such a data plane in four parts

- Hardware gadgets
  - Random number generators (RED, BLUE)
  - Binary tree of comparators (pFabric, SRPT)
- I/O interfaces
  - Drop/mark head/tail of queue
  - Interrupts for enqueue/dequeue
  - Rewrite packet fields
- State maintenance
  - Per-flow (WFQ, DRR)
  - Per-dst address (PF)
- A domain-specific instruction set
  - Expresses control flow
  - Implements new functions unavailable in hardware

Feasibility study: CoDel

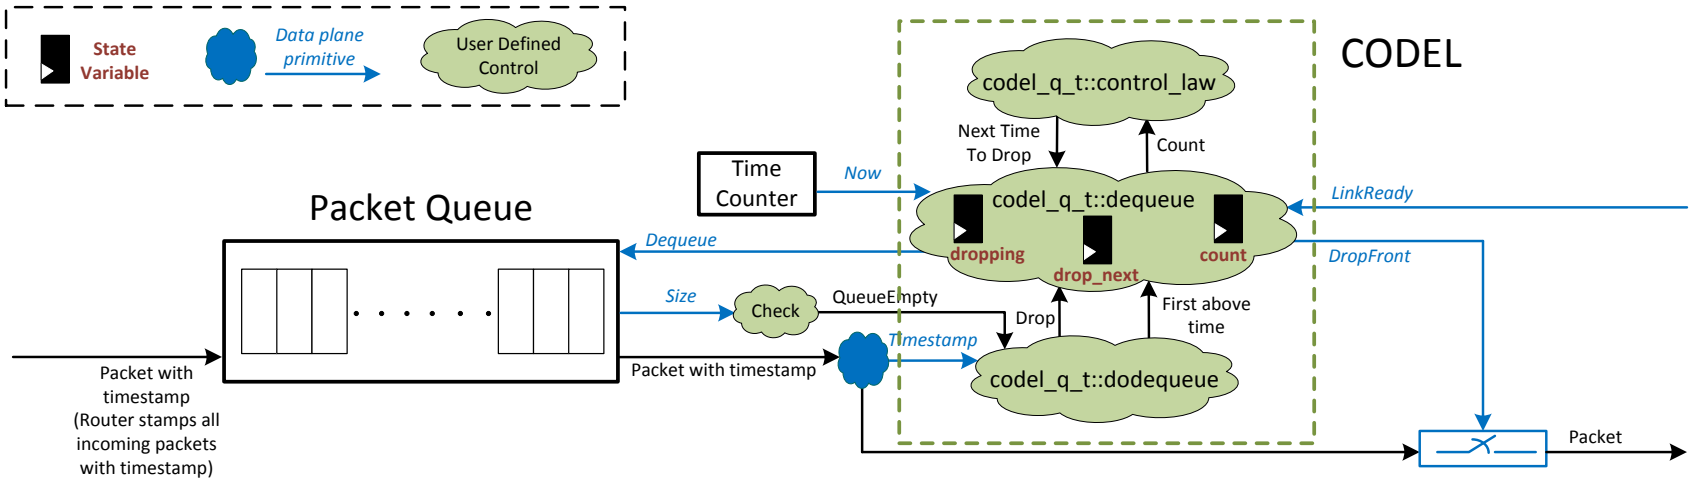

Synthesis numbers on the Xilinx Kintex-7

Resource	**Usage**	**Fraction**
Slice logic	1,256	1%
Slice logic dist.	1,975	2%
IO/GTX ports	27	2%
DSP slices	0	0%
Maximum speed	12.9 million pkts/s ~10 Gbps	

- Small fraction of the FPGA's resources.
- Can be improved by pipelining or parallelizing.

Limitations and Practical Considerations:

- Cannot express several network functions that need payloads.
- Mechanism to signal application objectives.
- Feasibility at 10G on high port-density switches.
- Energy, area, and performance costs of flexibility.

Related Work

- Active Networking, e.g., ANTS
- Software Routers, e.g., Click
- Software-Defined Networking, e.g., OpenFlow

Conclusion

- No silver bullet to in-network resource allocation.
- Algorithms will evolve: the data plane should help
- Reproduce our results: http://web.mit.edu/anirudh/www/sdn-data-plane.html